

Artificial Intelligence Nanodegree

Convolutional Neural Networks

Project: Write an Algorithm for a Dog Identification App

In this notebook, some template code has already been provided for you, and you will need to implement additional functionality to successfully complete this project. You will not need to modify the included code beyond what is requested. Sections that begin with '**(IMPLEMENTATION)**' in the header indicate that the following block of code will require additional functionality which you must provide. Instructions will be provided for each section, and the specifics of the implementation are marked in the code block with a 'TODO' statement. Please be sure to read the instructions carefully!

Note: Once you have completed all of the code implementations, you need to finalize your work by exporting the iPython Notebook as an HTML document. Before exporting the notebook to html, all of the code cells need to have been run so that reviewers can see the final implementation and output. You can then export the notebook by using the menu above and navigating to "\n", "**File -> Download as -> HTML (.html)**". Include the finished document along with this notebook as your submission.

In addition to implementing code, there will be questions that you must answer which relate to the project and your implementation. Each section where you will answer a question is preceded by a '**Question X**' header. Carefully read each question and provide thorough answers in the following text boxes that begin with '**Answer:**'. Your project submission will be evaluated based on your answers to each of the questions and the implementation you provide.


Note: Code and Markdown cells can be executed using the **Shift + Enter** keyboard shortcut. Markdown cells can be edited by double-clicking the cell to enter edit mode.

The rubric contains *optional* "Stand Out Suggestions" for enhancing the project beyond the minimum requirements. If you decide to pursue the "Stand Out Suggestions", you should include the code in this IPython notebook.

Why We're Here

In this notebook, you will make the first steps towards developing an algorithm that could be used as part of a mobile or web app. At the end of this project, your code will accept any user-supplied image as input. If a dog is detected in the image, it will provide an estimate of the dog's breed. If a human is detected, it will provide an estimate of the dog breed that is most resembling. The image below displays potential sample output of your finished project (... but we expect that each student's algorithm will behave differently!).

```
hello, dog!  
your predicted breed is ...  
American Staffordshire terrier
```



In this real-world setting, you will need to piece together a series of models to perform different tasks; for instance, the algorithm that detects humans in an image will be different from the CNN that infers dog breed. There are many points of possible failure, and no perfect algorithm exists. Your imperfect solution will nonetheless create a fun user experience!

The Road Ahead

We break the notebook into separate steps. Feel free to use the links below to navigate the notebook.

- [Step 0](#): Import Datasets
- [Step 1](#): Detect Humans
- [Step 2](#): Detect Dogs
- [Step 3](#): Create a CNN to Classify Dog Breeds (from Scratch)
- [Step 4](#): Use a CNN to Classify Dog Breeds (using Transfer Learning)
- [Step 5](#): Create a CNN to Classify Dog Breeds (using Transfer Learning)
- [Step 6](#): Write your Algorithm
- [Step 7](#): Test Your Algorithm

Step 0: Import Datasets

Import Dog Dataset

In the code cell below, we import a dataset of dog images. We populate a few variables through the use of the `load_files` function from the scikit-learn library:

- `train_files`, `valid_files`, `test_files` - numpy arrays containing file paths to images
- `train_targets`, `valid_targets`, `test_targets` - numpy arrays containing onehot-encoded classification labels
- `dog_names` - list of string-valued dog breed names for translating labels

```
In [1]: from sklearn.datasets import load_files
from keras.utils import np_utils
import numpy as np
from glob import glob

# define function to load train, test, and validation datasets
def load_dataset(path):
    data = load_files(path)
    dog_files = np.array(data['filenames'])
    dog_targets = np_utils.to_categorical(np.array(data['target']), 133)
    return dog_files, dog_targets

# load train, test, and validation datasets
train_files, train_targets = load_dataset('dogImages/train')
valid_files, valid_targets = load_dataset('dogImages/valid')
test_files, test_targets = load_dataset('dogImages/test')

# load list of dog names
dog_names = [item[20:-1] for item in
sorted(glob("dogImages/train/*/"))]

# print statistics about the dataset
print('There are %d total dog categories.' % len(dog_names))
print('There are %s total dog images.\n' %
len(np.hstack([train_files, valid_files, test_files])))
print('There are %d training dog images.' % len(train_files))
print('There are %d validation dog images.' % len(valid_files))
print('There are %d test dog images.' % len(test_files))
```

Using TensorFlow backend.

There are 133 total dog categories.
There are 8351 total dog images.

There are 6680 training dog images.
There are 835 validation dog images.
There are 836 test dog images.

Import Human Dataset

In the code cell below, we import a dataset of human images, where the file paths are stored in the numpy array `human_files`.

```
In [2]: import random
        random.seed(8675309)

        # load filenames in shuffled human dataset
        human_files = np.array(glob("lfw/**/*.jpg"))
        random.shuffle(human_files)

        # print statistics about the dataset
        print('There are %d total human images.' % len(human_files))
```

There are 13233 total human images.

Step 1: Detect Humans

We use OpenCV's implementation of Haar feature-based cascade classifiers (http://docs.opencv.org/trunk/d7/d8b/tutorial_py_face_detection.html) to detect human faces in images. OpenCV provides many pre-trained face detectors, stored as XML files on github (<https://github.com/opencv/opencv/tree/master/data/haarcascades>). We have downloaded one of these detectors and stored it in the `haarcascades` directory.

In the next code cell, we demonstrate how to use this detector to find human faces in a sample image.

```
In [3]: import cv2
import matplotlib.pyplot as plt
%matplotlib inline

# extract pre-trained face detector
face_cascade = cv2.CascadeClassifier('haarcascades/haarcascade_frontalface_alt.xml')

# load color (BGR) image
img = cv2.imread(human_files[3])

# convert BGR image to grayscale
gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)

# find faces in image
faces = face_cascade.detectMultiScale(gray)

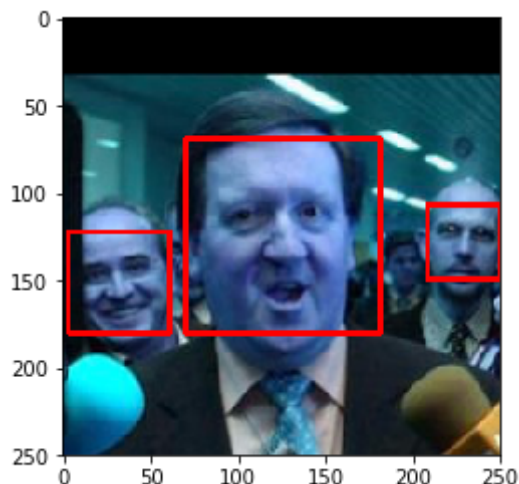
# print number of faces detected in the image
print('Number of faces detected:', len(faces))

# get bounding box for each detected face
for (x,y,w,h) in faces:
    # add bounding box to color image
    cv2.rectangle(img,(x,y),(x+w,y+h),(255,0,0),2)

# convert BGR image to RGB for plotting
cv_rgb = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)

# display the image, along with bounding box
plt.imshow(img)
plt.show()
```

Number of faces detected: 3



Before using any of the face detectors, it is standard procedure to convert the images to grayscale. The `detectMultiScale` function executes the classifier stored in `face_cascade` and takes the grayscale image as a parameter.

In the above code, `faces` is a numpy array of detected faces, where each row corresponds to a detected face. Each detected face is a 1D array with four entries that specifies the bounding box of the detected face. The first two entries in the array (extracted in the above code as `x` and `y`) specify the horizontal and vertical positions of the top left corner of the bounding box. The last two entries in the array (extracted here as `w` and `h`) specify the width and height of the box.

Write a Human Face Detector

We can use this procedure to write a function that returns `True` if a human face is detected in an image and `False` otherwise. This function, aptly named `face_detector`, takes a string-valued file path to an image as input and appears in the code block below.

```
In [4]: # returns "True" if face is detected in image stored at img_path
def face_detector(img_path):
    img = cv2.imread(img_path)
    gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)
    faces = face_cascade.detectMultiScale(gray)
    return len(faces) > 0
```

(IMPLEMENTATION) Assess the Human Face Detector

Question 1: Use the code cell below to test the performance of the `face_detector` function.

- What percentage of the first 100 images in `human_files` have a detected human face? 99%
- What percentage of the first 100 images in `dog_files` have a detected human face? 11%

Ideally, we would like 100% of human images with a detected face and 0% of dog images with a detected face. You will see that our algorithm falls short of this goal, but still gives acceptable performance. We extract the file paths for the first 100 images from each of the datasets and store them in the numpy arrays `human_files_short` and `dog_files_short`.

Answer:

Percentage of human faces detected in the human files: 99% Percentage of human faces detected in the dog files: 11%

```

In [5]: human_files_short = human_files[:100]
        dog_files_short = train_files[:100]
        # Do NOT modify the code above this line.

        ## TODO: Test the performance of the face_detector algorithm
        ## on the images in human_files_short and dog_files_short.
        detected_humans = 0
        detected_dogs = 0
        for human in human_files_short:
            if face_detector(human):
                detected_humans += 1
        for dog in dog_files_short:
            if face_detector(dog):
                detected_dogs += 1

        print(detected_humans)
        print(detected_dogs)

99
11

```

Question 2: This algorithmic choice necessitates that we communicate to the user that we accept human images only when they provide a clear view of a face (otherwise, we risk having unnecessarily frustrated users!). In your opinion, is this a reasonable expectation to pose on the user? If not, can you think of a way to detect humans in images that does not necessitate an image with a clearly presented face?

Answer: I don't think using images without humans or a clear view of a human faces are a problem. At the end we should provide more complete information to the user. Asking the user to only provide human faces could be a pain. Think about the case where only twenty images in a huge dataset of thousands of images are not humans. We should not stop our search by giving a message "please provide only human images". Probably the user is not able to tell where those wrong images are.

How could we use a method to detect humans in an image?

Let's think about what we know about how the human brain recognizes objects. Researchers at MIT's Department of Brain and Cognitive Sciences suggest that the human brain represents visual information in a hierarchical way. As visual input flows from the retina into primary visual cortex and then inferotemporal (IT) cortex, it is processed at each level and becomes more specific until objects can be identified. Well deep neural networks work in a similar fashion, so I think we definitely need a deep neural network with many layers to be able to filter the very specific features that define the shape of a human body.

We suggest the face detector from OpenCV as a potential way to detect human images in your algorithm, but you are free to explore other approaches, especially approaches that make use of deep learning :). Please use the code cell below to design and test your own face detection algorithm. If you decide to pursue this *optional* task, report performance on each of the datasets.

```

In [6]: ## (Optional) TODO: Report the performance of another
        ## face detection algorithm on the LFW dataset
        ### Feel free to use as many code cells as needed.

```


Step 2: Detect Dogs

In this section, we use a pre-trained [ResNet-50](https://ethereon.github.io/netscope/#/gist/db945b393d40bfa26006) ([http://ethereon.github.io/netscope/#/gist/db945b393d40bfa26006](https://ethereon.github.io/netscope/#/gist/db945b393d40bfa26006)) model to detect dogs in images. Our first line of code downloads the ResNet-50 model, along with weights that have been trained on [ImageNet](http://www.image-net.org/) (<http://www.image-net.org/>), a very large, very popular dataset used for image classification and other vision tasks. ImageNet contains over 10 million URLs, each linking to an image containing an object from one of [1000 categories](https://gist.github.com/yrevar/942d3a0ac09ec9e5eb3a) (<https://gist.github.com/yrevar/942d3a0ac09ec9e5eb3a>). Given an image, this pre-trained ResNet-50 model returns a prediction (derived from the available categories in ImageNet) for the object that is contained in the image.

```
In [6]: from keras.applications.resnet50 import ResNet50

# define ResNet50 model
ResNet50_model = ResNet50(weights='imagenet')
```

Pre-process the Data

When using TensorFlow as backend, Keras CNNs require a 4D array (which we'll also refer to as a 4D tensor) as input, with shape

$$(nb_samples, rows, columns, channels),$$

where `nb_samples` corresponds to the total number of images (or samples), and `rows`, `columns`, and `channels` correspond to the number of rows, columns, and channels for each image, respectively.

The `path_to_tensor` function below takes a string-valued file path to a color image as input and returns a 4D tensor suitable for supplying to a Keras CNN. The function first loads the image and resizes it to a square image that is 224×224 pixels. Next, the image is converted to an array, which is then resized to a 4D tensor. In this case, since we are working with color images, each image has three channels. Likewise, since we are processing a single image (or sample), the returned tensor will always have shape

$$(1, 224, 224, 3).$$

The `paths_to_tensor` function takes a numpy array of string-valued image paths as input and returns a 4D tensor with shape

$$(nb_samples, 224, 224, 3).$$

Here, `nb_samples` is the number of samples, or number of images, in the supplied array of image paths. It is best to think of `nb_samples` as the number of 3D tensors (where each 3D tensor corresponds to a different image) in your dataset!

```
In [7]: from keras.preprocessing import image
        from tqdm import tqdm

        def path_to_tensor(img_path):
            # loads RGB image as PIL.Image.Image type
            img = image.load_img(img_path, target_size=(224, 224))
            # convert PIL.Image.Image type to 3D tensor with shape (224, 224,
            3)
            x = image.img_to_array(img)
            # convert 3D tensor to 4D tensor with shape (1, 224, 224, 3) and
            return 4D tensor
            return np.expand_dims(x, axis=0)

        def paths_to_tensor(img_paths):
            list_of_tensors = [path_to_tensor(img_path) for img_path in
            tqdm(img_paths)]
            return np.vstack(list_of_tensors)
```

Making Predictions with ResNet-50

Getting the 4D tensor ready for ResNet-50, and for any other pre-trained model in Keras, requires some additional processing. First, the RGB image is converted to BGR by reordering the channels. All pre-trained models have the additional normalization step that the mean pixel (expressed in RGB as $[103.939, 116.779, 123.68]$ and calculated from all pixels in all images in ImageNet) must be subtracted from every pixel in each image. This is implemented in the imported function `preprocess_input`. If you're curious, you can check the code for `preprocess_input` [here](https://github.com/fchollet/keras/blob/master/keras/applications/imagenet_utils.py) (https://github.com/fchollet/keras/blob/master/keras/applications/imagenet_utils.py).

Now that we have a way to format our image for supplying to ResNet-50, we are now ready to use the model to extract the predictions. This is accomplished with the `predict` method, which returns an array whose i -th entry is the model's predicted probability that the image belongs to the i -th ImageNet category. This is implemented in the `ResNet50_predict_labels` function below.

By taking the argmax of the predicted probability vector, we obtain an integer corresponding to the model's predicted object class, which we can identify with an object category through the use of this [dictionary](https://gist.github.com/yrevar/942d3a0ac09ec9e5eb3a) (<https://gist.github.com/yrevar/942d3a0ac09ec9e5eb3a>).

```
In [8]: from keras.applications.resnet50 import preprocess_input, decode_predictions

        def ResNet50_predict_labels(img_path):
            # returns prediction vector for image located at img_path
            img = preprocess_input(path_to_tensor(img_path))
            return np.argmax(ResNet50_model.predict(img))
```

Write a Dog Detector

While looking at the [dictionary \(https://gist.github.com/yrevar/942d3a0ac09ec9e5eb3a\)](https://gist.github.com/yrevar/942d3a0ac09ec9e5eb3a), you will notice that the categories corresponding to dogs appear in an uninterrupted sequence and correspond to dictionary keys 151-268, inclusive, to include all categories from 'Chihuahua' to 'Mexican hairless'. Thus, in order to check to see if an image is predicted to contain a dog by the pre-trained ResNet-50 model, we need only check if the ResNet50_predict_labels function above returns a value between 151 and 268 (inclusive).

We use these ideas to complete the dog_detector function below, which returns True if a dog is detected in an image (and False if not).

```
In [9]: ### returns "True" if a dog is detected in the image stored at img_path
def dog_detector(img_path):
    prediction = ResNet50_predict_labels(img_path)
    return ((prediction <= 268) & (prediction >= 151))
```

(IMPLEMENTATION) Assess the Dog Detector

Question 3: Use the code cell below to test the performance of your dog_detector function.

- What percentage of the images in human_files_short have a detected dog?
- What percentage of the images in dog_files_short have a detected dog?

Answer: Dogs detected in human_files_short: 0% Dogs detected in dog_files_short: 99%

```
In [10]: ### TODO: Test the performance of the dog_detector function
### on the images in human_files_short and dog_files_short.

detected_humans = 0
detected_dogs = 0
for human in human_files_short:
    if dog_detector(human):
        detected_humans += 1
for dog in dog_files_short:
    if dog_detector(dog):
        detected_dogs += 1

print(detected_humans)
print(detected_dogs)

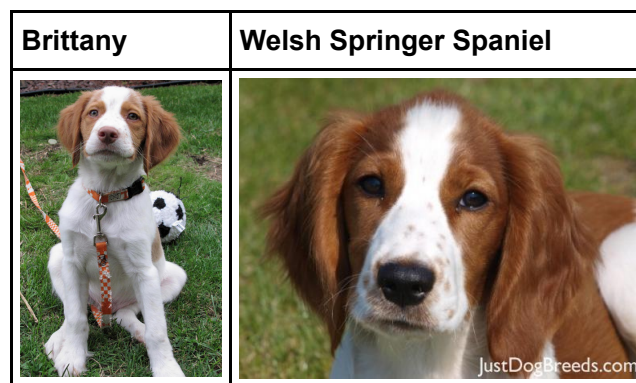
0
100
```

Step 3: Create a CNN to Classify Dog Breeds (from Scratch)

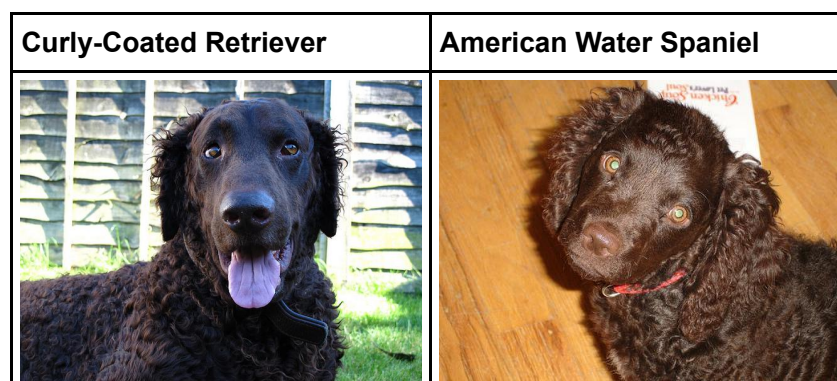
Now that we have functions for detecting humans and dogs in images, we need a way to predict breed from images. In this step, you will create a CNN that classifies dog breeds. You must create your CNN *from scratch* (so, you can't use transfer learning *yet!*), and you must attain a test accuracy of at least 1%. In Step 5 of this notebook, you will have the opportunity to use transfer learning to create a CNN that attains greatly improved accuracy.

Be careful with adding too many trainable layers! More parameters means longer training, which means you are more likely to need a GPU to accelerate the training process. Thankfully, Keras provides a handy estimate of the time that each epoch is likely to take; you can extrapolate this estimate to figure out how long it will take for your algorithm to train.

We mention that the task of assigning breed to dogs from images is considered exceptionally challenging. To see why, consider that *even a human* would have great difficulty in distinguishing between a Brittany and a Welsh Springer Spaniel.



It is not difficult to find other dog breed pairs with minimal inter-class variation (for instance, Curly-Coated Retrievers and American Water Spaniels).



Likewise, recall that labradors come in yellow, chocolate, and black. Your vision-based algorithm will have to conquer this high intra-class variation to determine how to classify all of these different shades as the same breed.

Yellow Labrador	Chocolate Labrador	Black Labrador
-----------------	--------------------	----------------



We also mention that random chance presents an exceptionally low bar: setting aside the fact that the classes are slightly imbalanced, a random guess will provide a correct answer roughly 1 in 133 times, which corresponds to an accuracy of less than 1%.

Remember that the practice is far ahead of the theory in deep learning. Experiment with many different architectures, and trust your intuition. And, of course, have fun!

Pre-process the Data

We rescale the images by dividing every pixel in every image by 255.

```
In [11]: from PIL import ImageFile
ImageFile.LOAD_TRUNCATED_IMAGES = True

# pre-process the data for Keras
train_tensors = paths_to_tensor(train_files).astype('float32')/255
valid_tensors = paths_to_tensor(valid_files).astype('float32')/255
test_tensors = paths_to_tensor(test_files).astype('float32')/255

100%|██████████| 6680/6680 [00:53<00:00, 125.66it/s]
100%|██████████| 835/835 [00:05<00:00, 140.57it/s]
100%|██████████| 836/836 [00:05<00:00, 141.57it/s]
```

(IMPLEMENTATION) Model Architecture

Create a CNN to classify dog breed. At the end of your code cell block, summarize the layers of your model by executing the line:

```
model.summary()
```

We have imported some Python modules to get you started, but feel free to import as many modules as you need. If you end up getting stuck, here's a hint that specifies a model that trains relatively fast on CPU and attains >1% test accuracy in 5 epochs:

Layer (type)	Output Shape	Param #
conv2d_1 (Conv2D)	(None, 223, 223, 16)	208
max_pooling2d_1 (MaxPooling2D)	(None, 111, 111, 16)	0
conv2d_2 (Conv2D)	(None, 110, 110, 32)	2080
max_pooling2d_2 (MaxPooling2D)	(None, 55, 55, 32)	0
conv2d_3 (Conv2D)	(None, 54, 54, 64)	8256
max_pooling2d_3 (MaxPooling2D)	(None, 27, 27, 64)	0
global_average_pooling2d_1 (GlobalAveragePooling2D)	(None, 64)	0
dense_1 (Dense)	(None, 133)	8645
Total params: 19,189.0		
Trainable params: 19,189.0		
Non-trainable params: 0.0		

INPUT

CONV

POOL

CONV

POOL

CONV

POOL

GAP

DENSE

Question 4: Outline the steps you took to get to your final CNN architecture and your reasoning at each step. If you chose to use the hinted architecture above, describe why you think that CNN architecture should work well for the image classification task.

Answer:

First we have to understand the problem we are trying to address. CNNs architecture will largely depend on the type and amount of data we have. In our case the problem is clear. We need to label the given image with its more probable dog breed and return an error if the image is neither a dog or a human. That would be a multiclass classification problem.

I based my architecture on the given CNN. It has a reasonable amount of convolutional layers that can get to more higher level features and it also does max pooling which is something I was planning to do. Pooling layers allow us to compute less and avoid overfitting. I trained it and it worked but didn't satisfy the requirements. So I started to look at previous examples, exercises and lectures in the course. I then added more layers: Three fully connected layers and two dropout layers in between to avoid overfitting and help against the vanishing gradient problem which makes training slow.

In one occasion I trained the model with the Global Average Pooling layer but it performed worst so I decided to delete that layer and add a Flatten layer instead.

I finally add a softmax function since we have a multiclass classification problem.

```
In [16]: from keras.layers import Conv2D, MaxPooling2D, GlobalAveragePooling2D
from keras.layers import Dropout, Flatten, Dense
from keras.models import Sequential

model = Sequential()

### TODO: Define your architecture.
model.add(Conv2D(filters=16, kernel_size=2, padding='same', activation='relu', input_shape=(224, 224, 3)))
model.add(MaxPooling2D(pool_size=2))
model.add(Conv2D(filters=32, kernel_size=2, padding='same', activation='relu'))
model.add(MaxPooling2D(pool_size=2))
model.add(Conv2D(filters=64, kernel_size=2, padding='same', activation='relu'))
model.add(MaxPooling2D(pool_size=2))
model.add(Dropout(0.3))
model.add(Flatten())
model.add(Dense(133, activation='relu'))
model.add(Dropout(0.4))
model.add(Dense(133, activation='softmax'))

model.summary()
```

Layer (type)	Output Shape	Param #
conv2d_4 (Conv2D)	(None, 224, 224, 16)	208
max_pooling2d_5 (MaxPooling2D)	(None, 112, 112, 16)	0
conv2d_5 (Conv2D)	(None, 112, 112, 32)	2080
max_pooling2d_6 (MaxPooling2D)	(None, 56, 56, 32)	0
conv2d_6 (Conv2D)	(None, 56, 56, 64)	8256
max_pooling2d_7 (MaxPooling2D)	(None, 28, 28, 64)	0
dropout_3 (Dropout)	(None, 28, 28, 64)	0
flatten_3 (Flatten)	(None, 50176)	0
dense_3 (Dense)	(None, 133)	6673541
dropout_4 (Dropout)	(None, 133)	0
dense_4 (Dense)	(None, 133)	17822
Total params: 6,701,907.0		
Trainable params: 6,701,907.0		
Non-trainable params: 0.0		

Compile the Model

```
In [17]: model.compile(optimizer='rmsprop', loss='categorical_crossentropy', metrics=['accuracy'])
```

(IMPLEMENTATION) Train the Model

Train your model in the code cell below. Use model checkpointing to save the model that attains the best validation loss.

You are welcome to [augment the training data \(https://blog.keras.io/building-powerful-image-classification-models-using-very-little-data.html\)](https://blog.keras.io/building-powerful-image-classification-models-using-very-little-data.html), but this is not a requirement.

```
In [18]: from keras.callbacks import ModelCheckpoint

        ### TODO: specify the number of epochs that you would like to use to
        train the model.

        epochs = 10

        ### Do NOT modify the code below this line.

        checkpointer = ModelCheckpoint(filepath='saved_models/weights.best.from_scratch.hdf5',
                                       verbose=1, save_best_only=True)

        model.fit(train_tensors, train_targets,
                  validation_data=(valid_tensors, valid_targets),
                  epochs=epochs, batch_size=20, callbacks=[checker], verbose=1)
```

Train on 6680 samples, validate on 835 samples

Epoch 1/10

6660/6680 [=====>.] - ETA: 0s - loss: 4.8911 - acc: 0.0131 Epoch 00000: val_loss improved from inf to 4.73451, saving model to saved_models/weights.best.from_scratch.hdf5

6680/6680 [=====] - 225s - loss: 4.8906 - acc: 0.0130 - val_loss: 4.7345 - val_acc: 0.0263

Epoch 2/10

6660/6680 [=====>.] - ETA: 0s - loss: 4.6051 - acc: 0.0353 Epoch 00001: val_loss improved from 4.73451 to 4.45931, saving model to saved_models/weights.best.from_scratch.hdf5

6680/6680 [=====] - 225s - loss: 4.6050 - acc: 0.0352 - val_loss: 4.4593 - val_acc: 0.0527

Epoch 3/10

6660/6680 [=====>.] - ETA: 0s - loss: 4.2851 - acc: 0.0644 Epoch 00002: val_loss improved from 4.45931 to 4.22461, saving model to saved_models/weights.best.from_scratch.hdf5

6680/6680 [=====] - 234s - loss: 4.2848 - acc: 0.0644 - val_loss: 4.2246 - val_acc: 0.0575

Epoch 4/10

6660/6680 [=====>.] - ETA: 0s - loss: 3.9524 - acc: 0.1080 Epoch 00003: val_loss improved from 4.22461 to 4.16529, saving model to saved_models/weights.best.from_scratch.hdf5

6680/6680 [=====] - 236s - loss: 3.9535 - acc: 0.1076 - val_loss: 4.1653 - val_acc: 0.0838

Epoch 5/10

6660/6680 [=====>.] - ETA: 0s - loss: 3.6069 - acc: 0.1560 Epoch 00004: val_loss did not improve

6680/6680 [=====] - 235s - loss: 3.6062 - acc: 0.1563 - val_loss: 4.2269 - val_acc: 0.0790

Epoch 6/10

6660/6680 [=====>.] - ETA: 0s - loss: 3.2333 - acc: 0.2135 Epoch 00005: val_loss did not improve

6680/6680 [=====] - 235s - loss: 3.2343 - acc: 0.2135 - val_loss: 4.2621 - val_acc: 0.0838

Epoch 7/10

6660/6680 [=====>.] - ETA: 0s - loss: 2.8618 - acc: 0.2907 Epoch 00006: val_loss did not improve

6680/6680 [=====] - 235s - loss: 2.8630 - acc: 0.2903 - val_loss: 4.3148 - val_acc: 0.0934

Epoch 8/10

6660/6680 [=====>.] - ETA: 0s - loss: 2.4716 - acc: 0.3746 Epoch 00007: val_loss did not improve

6680/6680 [=====] - 235s - loss: 2.4710 - acc: 0.3746 - val_loss: 4.6866 - val_acc: 0.0826

Epoch 9/10

6660/6680 [=====>.] - ETA: 0s - loss: 2.1729 - acc: 0.4461 Epoch 00008: val_loss did not improve

6680/6680 [=====] - 235s - loss: 2.1740 - acc: 0.4455 - val_loss: 4.7565 - val_acc: 0.0874

Epoch 10/10

6660/6680 [=====>.] - ETA: 0s - loss: 1.9215 - acc: 0.5038 Epoch 00009: val_loss did not improve

6680/6680 [=====] - 235s - loss: 1.9204 - acc: 0.5040 - val_loss: 5.3305 - val_acc: 0.0874

Out[18]: <keras.callbacks.History at 0x7f8920197f60>

Load the Model with the Best Validation Loss

```
In [19]: model.load_weights('saved_models/weights.best.from_scratch.hdf5')
```

Test the Model

Try out your model on the test dataset of dog images. Ensure that your test accuracy is greater than 1%.

```
In [20]: # get index of predicted dog breed for each image in test set
dog_breed_predictions = [np.argmax(model.predict(np.expand_dims(tensor, axis=0))) for tensor in test_tensors]

# report test accuracy
test_accuracy =
100*np.sum(np.array(dog_breed_predictions)==np.argmax(test_targets, axis=1))/len(dog_breed_predictions)
print('Test accuracy: %.4f%%' % test_accuracy)
```

Test accuracy: 10.2871%

Step 4: Use a CNN to Classify Dog Breeds

To reduce training time without sacrificing accuracy, we show you how to train a CNN using transfer learning. In the following step, you will get a chance to use transfer learning to train your own CNN.

Obtain Bottleneck Features

```
In [21]: bottleneck_features = np.load('bottleneck_features/DogVGG16Data.npz')
train_VGG16 = bottleneck_features['train']
valid_VGG16 = bottleneck_features['valid']
test_VGG16 = bottleneck_features['test']
```

Model Architecture

The model uses the the pre-trained VGG-16 model as a fixed feature extractor, where the last convolutional output of VGG-16 is fed as input to our model. We only add a global average pooling layer and a fully connected layer, where the latter contains one node for each dog category and is equipped with a softmax.

```
In [22]: VGG16_model = Sequential()
VGG16_model.add(GlobalAveragePooling2D(input_shape=train_VGG16.shape[1]
VGG16_model.add(Dense(133, activation='softmax'))
VGG16_model.summary()
```

Layer (type)	Output Shape	Param #
global_average_pooling2d_1 ((None, 512)	0
dense_5 (Dense)	(None, 133)	68229
Total params: 68,229.0		
Trainable params: 68,229.0		
Non-trainable params: 0.0		

Compile the Model

```
In [23]: VGG16_model.compile(loss='categorical_crossentropy', optimizer='rmsprop', metrics=['accuracy'])
```

Train the Model

```
In [24]: checkpointer = ModelCheckpoint(filepath='saved_models/weights.best.VG
G16.hdf5',
                                             verbose=1, save_best_only=True)

VGG16_model.fit(train_VGG16, train_targets,
                validation_data=(valid_VGG16, valid_targets),
                epochs=20, batch_size=20, callbacks=[checker],
                verbose=1)
```

Train on 6680 samples, validate on 835 samples

Epoch 1/20

6640/6680 [=====>.] - ETA: 0s - loss: 12.7361
- acc: 0.1142Epoch 00000: val_loss improved from inf to 11.62612, saving model to saved_models/weights.best.VGG16.hdf5

6680/6680 [=====] - 1s - loss: 12.7283 - acc: 0.1147 - val_loss: 11.6261 - val_acc: 0.1856

Epoch 2/20

6440/6680 [=====>..] - ETA: 0s - loss: 11.0035
- acc: 0.2410Epoch 00001: val_loss improved from 11.62612 to 10.79587, saving model to saved_models/weights.best.VGG16.hdf5

6680/6680 [=====] - 1s - loss: 10.9935 - acc: 0.2410 - val_loss: 10.7959 - val_acc: 0.2623

Epoch 3/20

6580/6680 [=====>.] - ETA: 0s - loss: 10.5836
- acc: 0.2903Epoch 00002: val_loss improved from 10.79587 to 10.74557, saving model to saved_models/weights.best.VGG16.hdf5

6680/6680 [=====] - 1s - loss: 10.5898 - acc: 0.2901 - val_loss: 10.7456 - val_acc: 0.2695

Epoch 4/20

6380/6680 [=====>..] - ETA: 0s - loss: 10.3092
- acc: 0.3204Epoch 00003: val_loss improved from 10.74557 to 10.49964, saving model to saved_models/weights.best.VGG16.hdf5

6680/6680 [=====] - 1s - loss: 10.3209 - acc: 0.3196 - val_loss: 10.4996 - val_acc: 0.2874

Epoch 5/20

6580/6680 [=====>.] - ETA: 0s - loss: 10.1633
- acc: 0.3372Epoch 00004: val_loss improved from 10.49964 to 10.43644, saving model to saved_models/weights.best.VGG16.hdf5

6680/6680 [=====] - 1s - loss: 10.1656 - acc: 0.3373 - val_loss: 10.4364 - val_acc: 0.2898

Epoch 6/20

6460/6680 [=====>.] - ETA: 0s - loss: 10.0739
- acc: 0.3483Epoch 00005: val_loss improved from 10.43644 to 10.32690, saving model to saved_models/weights.best.VGG16.hdf5

6680/6680 [=====] - 1s - loss: 10.0631 - acc: 0.3493 - val_loss: 10.3269 - val_acc: 0.3018

Epoch 7/20

6500/6680 [=====>.] - ETA: 0s - loss: 9.9058 - acc: 0.3646Epoch 00006: val_loss improved from 10.32690 to 10.16201, saving model to saved_models/weights.best.VGG16.hdf5

6680/6680 [=====] - 1s - loss: 9.8857 - acc: 0.3654 - val_loss: 10.1620 - val_acc: 0.3102

Epoch 8/20

6480/6680 [=====>.] - ETA: 0s - loss: 9.6864 - acc: 0.3772Epoch 00007: val_loss improved from 10.16201 to 9.98426, saving model to saved_models/weights.best.VGG16.hdf5

6680/6680 [=====] - 1s - loss: 9.6799 - acc: 0.3780 - val_loss: 9.9843 - val_acc: 0.3234

Epoch 9/20

6640/6680 [=====>.] - ETA: 0s - loss: 9.6382 - acc: 0.3889Epoch 00008: val_loss did not improve

6680/6680 [=====] - 1s - loss: 9.6217 - acc: 0.3898 - val_loss: 10.0783 - val_acc: 0.3210

Epoch 10/20

6420/6680 [=====>..] - ETA: 0s - loss: 9.5813 - acc: 0.3958Epoch 00009: val_loss did not improve

```
6680/6680 [=====] - 1s - loss: 9.5970 - acc:
0.3948 - val_loss: 10.0279 - val_acc: 0.3365
Epoch 11/20
6480/6680 [=====>.] - ETA: 0s - loss: 9.4896 -
acc: 0.3980Epoch 00010: val_loss improved from 9.98426 to 9.92733, sa
ving model to saved_models/weights.best.VGG16.hdf5
6680/6680 [=====] - 1s - loss: 9.5097 - acc:
0.3967 - val_loss: 9.9273 - val_acc: 0.3305
Epoch 12/20
6400/6680 [=====>..] - ETA: 0s - loss: 9.2655 -
acc: 0.4100Epoch 00011: val_loss improved from 9.92733 to 9.57406, sa
ving model to saved_models/weights.best.VGG16.hdf5
6680/6680 [=====] - 1s - loss: 9.2789 - acc:
0.4082 - val_loss: 9.5741 - val_acc: 0.3449
Epoch 13/20
6480/6680 [=====>.] - ETA: 0s - loss: 8.8098 -
acc: 0.4310Epoch 00012: val_loss improved from 9.57406 to 9.25239, sa
ving model to saved_models/weights.best.VGG16.hdf5
6680/6680 [=====] - 1s - loss: 8.8205 - acc:
0.4302 - val_loss: 9.2524 - val_acc: 0.3641
Epoch 14/20
6500/6680 [=====>.] - ETA: 0s - loss: 8.5618 -
acc: 0.4494Epoch 00013: val_loss improved from 9.25239 to 9.10297, sa
ving model to saved_models/weights.best.VGG16.hdf5
6680/6680 [=====] - 1s - loss: 8.5861 - acc:
0.4481 - val_loss: 9.1030 - val_acc: 0.3701
Epoch 15/20
6620/6680 [=====>.] - ETA: 0s - loss: 8.4821 -
acc: 0.4604Epoch 00014: val_loss improved from 9.10297 to 8.95623, sa
ving model to saved_models/weights.best.VGG16.hdf5
6680/6680 [=====] - 1s - loss: 8.4857 - acc:
0.4602 - val_loss: 8.9562 - val_acc: 0.3916
Epoch 16/20
6460/6680 [=====>.] - ETA: 0s - loss: 8.4527 -
acc: 0.4676Epoch 00015: val_loss did not improve
6680/6680 [=====] - 1s - loss: 8.4571 - acc:
0.4675 - val_loss: 9.0154 - val_acc: 0.3928
Epoch 17/20
6520/6680 [=====>.] - ETA: 0s - loss: 8.4392 -
acc: 0.4712Epoch 00016: val_loss did not improve
6680/6680 [=====] - 1s - loss: 8.4398 - acc:
0.4710 - val_loss: 8.9918 - val_acc: 0.3868
Epoch 18/20
6460/6680 [=====>.] - ETA: 0s - loss: 8.3969 -
acc: 0.4729Epoch 00017: val_loss did not improve
6680/6680 [=====] - 1s - loss: 8.4183 - acc:
0.4716 - val_loss: 8.9660 - val_acc: 0.3916
Epoch 19/20
6640/6680 [=====>.] - ETA: 0s - loss: 8.3800 -
acc: 0.4744Epoch 00018: val_loss improved from 8.95623 to 8.93128, sa
ving model to saved_models/weights.best.VGG16.hdf5
6680/6680 [=====] - 1s - loss: 8.3757 - acc:
0.4747 - val_loss: 8.9313 - val_acc: 0.3940
Epoch 20/20
6560/6680 [=====>.] - ETA: 0s - loss: 8.3658 -
acc: 0.4768Epoch 00019: val_loss did not improve
```



```
6680/6680 [=====] - 1s - loss: 8.3523 - acc:
0.4777 - val_loss: 8.9343 - val_acc: 0.3916
```

```
Out[24]: <keras.callbacks.History at 0x7f882c3f2b70>
```

Load the Model with the Best Validation Loss

```
In [25]: VGG16_model.load_weights('saved_models/weights.best.VGG16.hdf5')
```

Test the Model

Now, we can use the CNN to test how well it identifies breed within our test dataset of dog images. We print the test accuracy below.

```
In [26]: # get index of predicted dog breed for each image in test set
VGG16_predictions = [np.argmax(VGG16_model.predict(np.expand_dims(feature, axis=0))) for feature in test_VGG16]

# report test accuracy
test_accuracy = 100*np.sum(np.array(VGG16_predictions)==np.argmax(test_targets, axis=1))/len(VGG16_predictions)
print('Test accuracy: %.4f%%' % test_accuracy)
```

```
Test accuracy: 37.9187%
```

Predict Dog Breed with the Model

```
In [27]: from extract_bottleneck_features import *

def VGG16_predict_breed(img_path):
    # extract bottleneck features
    bottleneck_feature = extract_VGG16(path_to_tensor(img_path))
    # obtain predicted vector
    predicted_vector = VGG16_model.predict(bottleneck_feature)
    # return dog breed that is predicted by the model
    return dog_names[np.argmax(predicted_vector)]
```

Step 5: Create a CNN to Classify Dog Breeds (using Transfer Learning)

You will now use transfer learning to create a CNN that can identify dog breed from images. Your CNN must attain at least 60% accuracy on the test set.

In Step 4, we used transfer learning to create a CNN using VGG-16 bottleneck features. In this section, you must use the bottleneck features from a different pre-trained model. To make things easier for you, we have pre-computed the features for all of the networks that are currently available in Keras:

- [VGG-19 \(https://s3-us-west-1.amazonaws.com/udacity-aind/dog-project/DogVGG19Data.npz\)](https://s3-us-west-1.amazonaws.com/udacity-aind/dog-project/DogVGG19Data.npz) bottleneck features
- [ResNet-50 \(https://s3-us-west-1.amazonaws.com/udacity-aind/dog-project/DogResnet50Data.npz\)](https://s3-us-west-1.amazonaws.com/udacity-aind/dog-project/DogResnet50Data.npz) bottleneck features
- [Inception \(https://s3-us-west-1.amazonaws.com/udacity-aind/dog-project/DogInceptionV3Data.npz\)](https://s3-us-west-1.amazonaws.com/udacity-aind/dog-project/DogInceptionV3Data.npz) bottleneck features
- [Xception \(https://s3-us-west-1.amazonaws.com/udacity-aind/dog-project/DogXceptionData.npz\)](https://s3-us-west-1.amazonaws.com/udacity-aind/dog-project/DogXceptionData.npz) bottleneck features

The files are encoded as such:

```
Dog{network}Data.npz
```

where {network}, in the above filename, can be one of VGG19, Resnet50, InceptionV3, or Xception. Pick one of the above architectures, download the corresponding bottleneck features, and store the downloaded file in the `bottleneck_features/` folder in the repository.

(IMPLEMENTATION) Obtain Bottleneck Features

In the code block below, extract the bottleneck features corresponding to the train, test, and validation sets by running the following:

```
bottleneck_features = np.load('bottleneck_features/Dog{network}Data.npz')
train_{network} = bottleneck_features['train']
valid_{network} = bottleneck_features['valid']
test_{network} = bottleneck_features['test']
```

```
In [28]: ### TODO: Obtain bottleneck features from another pre-trained CNN.
# bottleneck_features = np.load('bottleneck_features/DogVGG19Data.npz')
# bottleneck_features = np.load('bottleneck_features/DogResnet50Data.npz')
# bottleneck_features = np.load('bottleneck_features/DogInceptionV3Data.npz')
bottleneck_features = np.load('bottleneck_features/DogXceptionData.npz')
breed_train = bottleneck_features['train']
breed_valid = bottleneck_features['valid']
breed_test = bottleneck_features['test']
```

(IMPLEMENTATION) Model Architecture

Create a CNN to classify dog breed. At the end of your code cell block, summarize the layers of your model by executing the line:

```
<your model's name>.summary()
```

Question 5: Outline the steps you took to get to your final CNN architecture and your reasoning at each step. Describe why you think the architecture is suitable for the current problem.

Answer:

Adding a Global Average Pooling in this case is enough to reduce the amount of parameters. A fully connected layer with softmax will also be enough to get the probability array.

```
In [29]: ### TODO: Define your architecture.
breed_model = Sequential()
breed_model.add(GlobalAveragePooling2D(input_shape=breed_train.shape[1:]

breed_model.add(Dense(133, activation='softmax'))

breed_model.summary()
```

Layer (type)	Output Shape	Param #
global_average_pooling2d_2 ((None, 2048)		0
dense_6 (Dense)	(None, 133)	272517
Total params: 272,517.0		
Trainable params: 272,517.0		
Non-trainable params: 0.0		

(IMPLEMENTATION) Compile the Model

```
In [30]: ### TODO: Compile the model.  
breed_model.compile(loss='categorical_crossentropy', optimizer='rmsprop', metrics=['accuracy'])
```

(IMPLEMENTATION) Train the Model

Train your model in the code cell below. Use model checkpointing to save the model that attains the best validation loss.

You are welcome to [augment the training data \(https://blog.keras.io/building-powerful-image-classification-models-using-very-little-data.html\)](https://blog.keras.io/building-powerful-image-classification-models-using-very-little-data.html), but this is not a requirement.

```
In [40]: ### TODO: Train the model.
# checkpointer = ModelCheckpoint(filepath='saved_models/weights.best.
VGG16.hdf5',
#                                     verbose=1, save_best_only=True)
# checkpointer = ModelCheckpoint(filepath='saved_models/weights.best.
VGG19.hdf5',
#                                     verbose=1, save_best_only=True)
# checkpointer = ModelCheckpoint(filepath='saved_models/weights.best.
Resnet50.hdf5',
#                                     verbose=1, save_best_only=True)
# checkpointer = ModelCheckpoint(filepath='saved_models/weights.best.
InceptionV3.hdf5',
#                                     verbose=1, save_best_only=True)
checker = ModelCheckpoint(filepath='saved_models/weights.best.Xc
eption.hdf5',
                           verbose=1, save_best_only=True)

breed_model.fit(breed_train, train_targets,
                validation_data=(breed_valid, valid_targets),
                epochs=20, batch_size=20, callbacks=[checker],
                verbose=1)
```

Train on 6680 samples, validate on 835 samples

Epoch 1/20

6620/6680 [=====>.] - ETA: 0s - loss: 0.2253 - acc: 0.9304 Epoch 00000: val_loss improved from inf to 0.52316, saving model to saved_models/weights.best.Xception.hdf5

6680/6680 [=====] - 3s - loss: 0.2250 - acc: 0.9305 - val_loss: 0.5232 - val_acc: 0.8647

Epoch 2/20

6620/6680 [=====>.] - ETA: 0s - loss: 0.1938 - acc: 0.9363 Epoch 00001: val_loss improved from 0.52316 to 0.51183, saving model to saved_models/weights.best.Xception.hdf5

6680/6680 [=====] - 3s - loss: 0.1933 - acc: 0.9364 - val_loss: 0.5118 - val_acc: 0.8611

Epoch 3/20

6620/6680 [=====>.] - ETA: 0s - loss: 0.1761 - acc: 0.9431 Epoch 00002: val_loss did not improve

6680/6680 [=====] - 3s - loss: 0.1754 - acc: 0.9431 - val_loss: 0.5727 - val_acc: 0.8491

Epoch 4/20

6640/6680 [=====>.] - ETA: 0s - loss: 0.1606 - acc: 0.9492 Epoch 00003: val_loss did not improve

6680/6680 [=====] - 3s - loss: 0.1613 - acc: 0.9491 - val_loss: 0.5848 - val_acc: 0.8503

Epoch 5/20

6620/6680 [=====>.] - ETA: 0s - loss: 0.1481 - acc: 0.9548 Epoch 00004: val_loss did not improve

6680/6680 [=====] - 3s - loss: 0.1480 - acc: 0.9548 - val_loss: 0.5709 - val_acc: 0.8527

Epoch 6/20

6620/6680 [=====>.] - ETA: 0s - loss: 0.1344 - acc: 0.9566 Epoch 00005: val_loss did not improve

6680/6680 [=====] - 3s - loss: 0.1347 - acc: 0.9567 - val_loss: 0.6015 - val_acc: 0.8491

Epoch 7/20

6620/6680 [=====>.] - ETA: 0s - loss: 0.1201 - acc: 0.9630 Epoch 00006: val_loss did not improve

6680/6680 [=====] - 3s - loss: 0.1197 - acc: 0.9630 - val_loss: 0.5777 - val_acc: 0.8623

Epoch 8/20

6620/6680 [=====>.] - ETA: 0s - loss: 0.1126 - acc: 0.9660 Epoch 00007: val_loss did not improve

6680/6680 [=====] - 3s - loss: 0.1135 - acc: 0.9659 - val_loss: 0.6008 - val_acc: 0.8539

Epoch 9/20

6620/6680 [=====>.] - ETA: 0s - loss: 0.1035 - acc: 0.9696 Epoch 00008: val_loss did not improve

6680/6680 [=====] - 3s - loss: 0.1028 - acc: 0.9699 - val_loss: 0.6421 - val_acc: 0.8491

Epoch 10/20

6600/6680 [=====>.] - ETA: 0s - loss: 0.0947 - acc: 0.9739 Epoch 00009: val_loss did not improve

6680/6680 [=====] - 3s - loss: 0.0946 - acc: 0.9738 - val_loss: 0.6959 - val_acc: 0.8515

Epoch 11/20

6640/6680 [=====>.] - ETA: 0s - loss: 0.0897 - acc: 0.9736 Epoch 00010: val_loss did not improve

6680/6680 [=====] - 3s - loss: 0.0898 - acc:

```

0.9735 - val_loss: 0.6428 - val_acc: 0.8467
Epoch 12/20
6620/6680 [=====>.] - ETA: 0s - loss: 0.0843 -
acc: 0.9760Epoch 00011: val_loss did not improve
6680/6680 [=====] - 3s - loss: 0.0864 - acc:
0.9759 - val_loss: 0.6696 - val_acc: 0.8611
Epoch 13/20
6580/6680 [=====>.] - ETA: 0s - loss: 0.0810 -
acc: 0.9764Epoch 00012: val_loss did not improve
6680/6680 [=====] - 3s - loss: 0.0805 - acc:
0.9766 - val_loss: 0.6743 - val_acc: 0.8491
Epoch 14/20
6640/6680 [=====>.] - ETA: 0s - loss: 0.0762 -
acc: 0.9789Epoch 00013: val_loss did not improve
6680/6680 [=====] - 3s - loss: 0.0760 - acc:
0.9790 - val_loss: 0.6738 - val_acc: 0.8575
Epoch 15/20
6620/6680 [=====>.] - ETA: 0s - loss: 0.0698 -
acc: 0.9801Epoch 00014: val_loss did not improve
6680/6680 [=====] - 3s - loss: 0.0693 - acc:
0.9802 - val_loss: 0.7002 - val_acc: 0.8539
Epoch 16/20
6620/6680 [=====>.] - ETA: 0s - loss: 0.0678 -
acc: 0.9804Epoch 00015: val_loss did not improve
6680/6680 [=====] - 3s - loss: 0.0679 - acc:
0.9802 - val_loss: 0.7103 - val_acc: 0.8599
Epoch 17/20
6620/6680 [=====>.] - ETA: 0s - loss: 0.0611 -
acc: 0.9819Epoch 00016: val_loss did not improve
6680/6680 [=====] - 3s - loss: 0.0608 - acc:
0.9819 - val_loss: 0.7391 - val_acc: 0.8467
Epoch 18/20
6560/6680 [=====>.] - ETA: 0s - loss: 0.0605 -
acc: 0.9816Epoch 00017: val_loss did not improve
6680/6680 [=====] - 3s - loss: 0.0607 - acc:
0.9816 - val_loss: 0.7060 - val_acc: 0.8515
Epoch 19/20
6580/6680 [=====>.] - ETA: 0s - loss: 0.0564 -
acc: 0.9840Epoch 00018: val_loss did not improve
6680/6680 [=====] - 3s - loss: 0.0556 - acc:
0.9843 - val_loss: 0.7433 - val_acc: 0.8611
Epoch 20/20
6620/6680 [=====>.] - ETA: 0s - loss: 0.0529 -
acc: 0.9850Epoch 00019: val_loss did not improve
6680/6680 [=====] - 3s - loss: 0.0526 - acc:
0.9850 - val_loss: 0.7291 - val_acc: 0.8647

```

```
Out[40]: <keras.callbacks.History at 0x7f8994664c18>
```

(IMPLEMENTATION) Load the Model with the Best Validation Loss

```
In [41]: ### TODO: Load the model weights with the best validation loss.
# breed_model.load_weights('saved_models/weights.best.VGG19.hdf5')
# breed_model.load_weights('saved_models/weights.best.Resnet50.hdf5')
# breed_model.load_weights('saved_models/weights.best.InceptionV3.hdf5')
breed_model.load_weights('saved_models/weights.best.Xception.hdf5')
```

(IMPLEMENTATION) Test the Model

Try out your model on the test dataset of dog images. Ensure that your test accuracy is greater than 60%.

```
In [42]: ### TODO: Calculate classification accuracy on the test dataset.
# get index of predicted dog breed for each image in test set
breed_predictions = [np.argmax(breed_model.predict(np.expand_dims(feature, axis=0))) for feature in breed_test]

# report test accuracy
test_accuracy = 100*np.sum(np.array(breed_predictions)==np.argmax(test_targets, axis=1))/len(breed_predictions)
print('Test accuracy: %.4f%%' % test_accuracy)
```

Test accuracy: 85.6459%

(IMPLEMENTATION) Predict Dog Breed with the Model

Write a function that takes an image path as input and returns the dog breed (Affenpinscher, Afghan_hound, etc) that is predicted by your model.

Similar to the analogous function in Step 5, your function should have three steps:

1. Extract the bottleneck features corresponding to the chosen CNN model.
2. Supply the bottleneck features as input to the model to return the predicted vector. Note that the `argmax` of this prediction vector gives the index of the predicted dog breed.
3. Use the `dog_names` array defined in Step 0 of this notebook to return the corresponding breed.

The functions to extract the bottleneck features can be found in `extract_bottleneck_features.py`, and they have been imported in an earlier code cell. To obtain the bottleneck features corresponding to your chosen CNN architecture, you need to use the function

```
extract_{network}
```

where `{network}`, in the above filename, should be one of VGG19, Resnet50, InceptionV3, or Xception.


```
In [43]: ### TODO: Write a function that takes a path to an image as input  
### and returns the dog breed that is predicted by the model.  
from extract_bottleneck_features import *  
  
extractors = {  
    'VGG16': extract_VGG16,  
    'VGG19': extract_VGG19,  
    'Resnet50': extract_Resnet50,  
    'Xception': extract_Xception,  
    'InceptionV3': extract_InceptionV3  
}  
  
def predict_breed(img_path, bf_name):  
    # extract bottleneck features  
    bottleneck_feature = extractors[bf_name]  
(path_to_tensor(img_path))  
    # obtain predicted vector  
    predicted_vector = breed_model.predict(bottleneck_feature)  
    # return dog breed that is predicted by the model  
    return dog_names[np.argmax(predicted_vector)]
```


Step 6: Write your Algorithm

Write an algorithm that accepts a file path to an image and first determines whether the image contains a human, dog, or neither. Then,

- if a **dog** is detected in the image, return the predicted breed.
- if a **human** is detected in the image, return the resembling dog breed.
- if **neither** is detected in the image, provide output that indicates an error.

You are welcome to write your own functions for detecting humans and dogs in images, but feel free to use the `face_detector` and `dog_detector` functions developed above. You are **required** to use your CNN from Step 5 to predict dog breed.

Some sample output for our algorithm is provided below, but feel free to design your own user experience!



(IMPLEMENTATION) Write your Algorithm

```
In [44]: ### TODO: Write your algorithm.  
### Feel free to use as many code cells as needed.  
bf_name = 'Xception'  
  
def prediction_machine(img_path):  
    if face_detector(img_path):  
        print('Hello, human!')  
        prediction = predict_breed(img_path, bf_name)  
    elif dog_detector(img_path):  
        print('Hello, dog!')  
        prediction = predict_breed(img_path, bf_name)  
    else:  
        return False  
  
    return prediction
```

Step 7: Test Your Algorithm

In this section, you will take your new algorithm for a spin! What kind of dog does the algorithm think that **you** look like? If you have a dog, does it predict your dog's breed accurately? If you have a cat, does it mistakenly think that your cat is a dog?

(IMPLEMENTATION) Test Your Algorithm on Sample Images!

Test your algorithm at least six images on your computer. Feel free to use any images you like. Use at least two human and two dog images.

Question 6: Is the output better than you expected :) ? Or worse :(? Provide at least three possible points of improvement for your algorithm.

Answer:


Output is not that bad. It is not perfect. I first picked a random image from my personal dataset of images. But I wanted to be able to see how all my images were labeled at once, so I decided to add a loop to my algorithm and print all labels.

Talking about improvements I think speed is one for sure, it takes a lot of time to compute all labels for my images and it is super boring to wait. UI and UX should be improved in the next version.

```
In [45]: import matplotlib.pyplot as plt
import numpy as np
from glob import glob
import cv2

def get_breed():
    personal_images = np.array(glob("images/personal/*"))
    for img_path in personal_images:
        p_img = cv2.imread(img_path)
        prediction = prediction_machine(img_path)
        plt.imshow(p_img)
        plt.show()
        if prediction is False:
            print('ERROR: The input image does not correspond to a do
g nor a human.')
        else:
            print('You look like a:
{0}'.format(prediction.replace('_', ' ')))
            print('=====')
            print('====='')
            print('\n')
            print('\n')
get_breed()
```


Hello, dog!



You look like a: Curly-coated retriever

=====

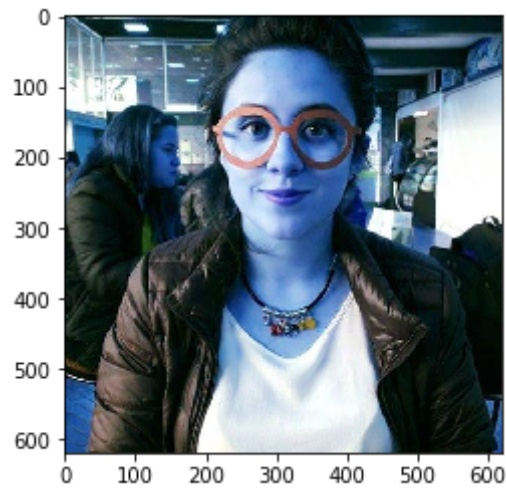
Hello, dog!



You look like a: Dogue de bordeaux

=====

Hello, human!



You look like a: Dachshund

=====


Hello, human!



You look like a: Chinese crested

=====


Hello, human!



You look like a: Greyhound

=====


Hello, dog!



You look like a: Labrador retriever

=====

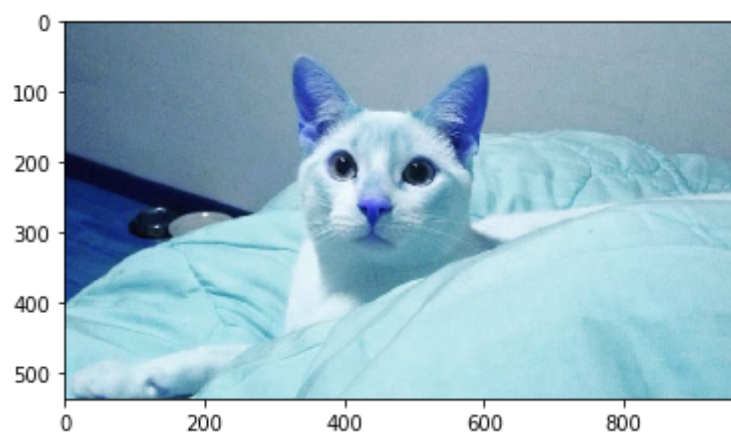
Hello, dog!



You look like a: Labrador retriever

=====

Hello, human!



You look like a: Chihuahua

=====


Hello, human!



You look like a: Dachshund

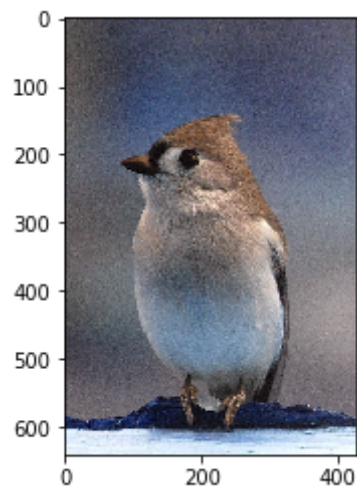
=====

Hello, dog!




You look like a: Labrador retriever

=====



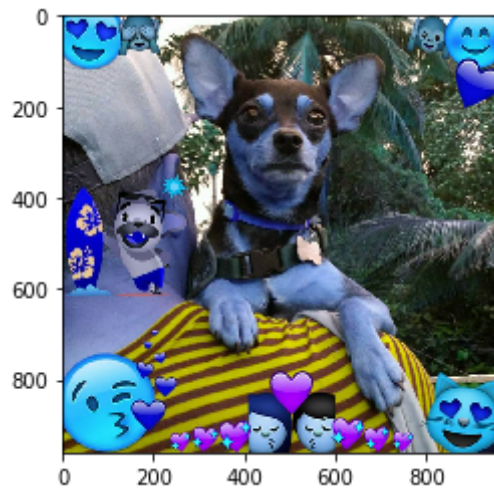
ERROR: The input image does not correspond to a dog nor a human.

Hello, dog!



You look like a: Curly-coated retriever


Hello, dog!



You look like a: Chihuahua

=====

Hello, dog!



You look like a: Manchester terrier

=====

In []: